

OLUMIDE ERINFOLAMI

Machine Learning Engineer, Applied ML for Fintech & Product Teams

Lagos, Nigeria (open to hybrid) · anthonyerinfoami@gmail.com · [linkedin.com/in/olumide-erinfoami](https://www.linkedin.com/in/olumide-erinfoami) · github.com/Mocdew
· olumide-erinfoami.vercel.app

SUMMARY

Computer Science graduate and former ML intern who ships applied machine learning as production-ready software, a payment-retry engine that recovered 31% more revenue in simulation, and a loan-default risk model trained on 255K loans. Strong in tabular ML, forecasting, and explainable AI; expanding into NLP and LLM applications.

TECHNICAL SKILLS

Languages: Python, SQL

Machine Learning: scikit-learn, LightGBM, XGBoost, CatBoost, Optuna, SHAP, model calibration, cost-sensitive learning

Deep Learning & NLP: PyTorch (LSTM, GRU), TF-IDF, embeddings, text classification

Forecasting & Statistics: time-series (ARIMA, GARCH), Bayesian modeling, multi-armed bandits, A/B experimentation

MLOps & Serving: Flask REST APIs, pytest, model versioning, GitHub Actions (CI/CD), Streamlit, Gradio

Data: pandas, NumPy, Plotly, Jupyter

EXPERIENCE

Machine Learning Intern, Notzero Technologies

Mar 2026 – Sep 2026

Lagos, Nigeria (Hybrid)

- Recovered **31% more revenue per failed payment** in simulation (\$15.93 vs. \$12.12 baseline) by building Recoup, an ML system that decides when and how often to retry a failed subscription charge, replacing the industry-standard fixed “retry after 1 / 3 / 7 days” rule.
- Drove **20–35% revenue gains across six stress-test scenarios** (and **+19% for small merchants**) by modeling each customer’s probability of paying at every retry window with probabilistic models in Python and SciPy, sharing signal across merchants so low-history businesses benefit.
- Safeguarded rollout with a **statistical safety gate** that blocks deployment until real-world evidence confirms the lift, so the business never bets revenue on an unproven model.

PROJECTS

Can I Borrow.ai, Loan Default Risk Engine · *LightGBM, scikit-learn, SHAP, Flask, pytest*

- Trained an end-to-end default-prediction system on **255K consumer loans**, reaching **ROC-AUC 0.76** on 51K unseen loans with well-calibrated risk (predicted 11.70% vs. actual 11.61% default rate) and a cost-tuned threshold that catches **80% of defaults**.
- Engineered for trust and auditability: SHAP explanations for every decision, protected attributes excluded, automated fairness checks on each run, and versioned models with auto-generated model cards, served through a Flask API and web app.

Financial Time-Series Forecasting, Oil, Inflation & FX · *statsmodels, LightGBM, PyTorch*

- Forecast oil-price volatility across **39 years of data at 4.6% error vs. 60% for a classical GARCH baseline**; built LSTM and GRU models in PyTorch for inflation and currency forecasting.

DSN 2026 Qualification Hackathon, 27th of 250 (top 11%) · *Kaggle, LightGBM, XGBoost, CatBoost, CI*

- Showed a simple, explainable store-level model (RMSE 1,070.8) **beat tuned LightGBM, XGBoost and CatBoost** (1,097–1,101) under repeated cross-validation; cut error on never-seen stores by **18%** and shipped calibrated 80% prediction intervals (79.7% actual coverage) as a tested Python package with GitHub Actions CI.

CrediPulse, Explainable Credit Scoring Platform · *LightGBM, Optuna, SHAP, Flask, JavaScript*

- Tuned a LightGBM credit model (**ROC-AUC 0.81**) behind a web app that returns a 300–850 score, an approve/decline decision, the reasons behind it, and bulk CSV scoring.

Deal Desk, Creator Brand-Deal Platform (MVP) · *Flask, Google OAuth, YouTube API, pytest*

- Built a SaaS MVP that connects a creator’s YouTube account, tracks sponsored-post performance against their baseline, and generates shareable PDF reports for brands.

EDUCATION

Covenant University, B.Sc. Computer Science

2021 – 2025

Ota, Nigeria