

# OLUMIDE ERINFOLAMI

## Machine Learning Engineer — Applied ML for Fintech & Product Teams

Lagos, Nigeria (open to hybrid) · [anthonyerinfoami@gmail.com](mailto:anthonyerinfoami@gmail.com) · [linkedin.com/in/olumide-erinfoami](https://linkedin.com/in/olumide-erinfoami) · [github.com/Mocdew](https://github.com/Mocdew)  
· [olumide-erinfoami.vercel.app](https://olumide-erinfoami.vercel.app)

## SUMMARY

Computer Science graduate and former ML intern who ships applied machine learning as production-ready software — a payment-retry engine that recovered 31% more revenue in simulation, and a loan-default risk model trained on 255K loans. Strong in tabular ML, forecasting, and explainable AI; expanding into NLP and LLM applications.

## TECHNICAL SKILLS

**Languages:** Python, SQL

**Machine Learning:** scikit-learn, LightGBM, XGBoost, CatBoost, Optuna, SHAP, model calibration, cost-sensitive learning

**Deep Learning & NLP:** PyTorch (LSTM, GRU), TF-IDF, embeddings, text classification

**Forecasting & Statistics:** time-series (ARIMA, GARCH), Bayesian modeling, multi-armed bandits, A/B experimentation

**MLOps & Serving:** Flask REST APIs, pytest, model versioning, GitHub Actions (CI/CD), Streamlit, Gradio

**Data:** pandas, NumPy, Plotly, Jupyter

## EXPERIENCE

### Machine Learning Intern — Notzero Technologies

Mar 2026 – Sep 2026

Lagos, Nigeria (Hybrid)

- Recovered **31% more revenue per failed payment** in simulation (\$15.93 vs. \$12.12 baseline) by building Recoup, an ML system that decides when and how often to retry a failed subscription charge — replacing the industry-standard fixed “retry after 1 / 3 / 7 days” rule.
- Drove **20–35% revenue gains across six stress-test scenarios** (and **+19% for small merchants**) by modeling each customer’s probability of paying at every retry window with probabilistic models in Python and SciPy, sharing signal across merchants so low-history businesses benefit.
- Safeguarded rollout with a **statistical safety gate** that blocks deployment until real-world evidence confirms the lift, so the business never bets revenue on an unproven model.

## PROJECTS

### Can I Borrow.ai — Loan Default Risk Engine · *LightGBM, scikit-learn, SHAP, Flask, pytest*

- Trained an end-to-end default-prediction system on **255K consumer loans**, reaching **ROC-AUC 0.76** on 51K unseen loans with well-calibrated risk (predicted 11.70% vs. actual 11.61% default rate) and a cost-tuned threshold that catches **80% of defaults**.
- Engineered for trust and auditability: SHAP explanations for every decision, protected attributes excluded, automated fairness checks on each run, and versioned models with auto-generated model cards — served through a Flask API and web app.

### Financial Time-Series Forecasting — Oil, Inflation & FX · *statsmodels, LightGBM, PyTorch*

- Forecast oil-price volatility across **39 years of data at 4.6% error vs. 60% for a classical GARCH baseline**; built LSTM and GRU models in PyTorch for inflation and currency forecasting.

### DSN 2026 Qualification Hackathon — 27th of 250 (top 11%) · *Kaggle, LightGBM, XGBoost, CatBoost, CI*

- Showed a simple, explainable store-level model (RMSE 1,070.8) **beat tuned LightGBM, XGBoost and CatBoost** (1,097–1,101) under repeated cross-validation; cut error on never-seen stores by **18%** and shipped calibrated 80% prediction intervals (79.7% actual coverage) as a tested Python package with GitHub Actions CI.

### CrediPulse — Explainable Credit Scoring Platform · *LightGBM, Optuna, SHAP, Flask, JavaScript*

- Tuned a LightGBM credit model (**ROC-AUC 0.81**) behind a web app that returns a 300–850 score, an approve/decline decision, the reasons behind it, and bulk CSV scoring.

### Deal Desk — Creator Brand-Deal Platform (MVP) · *Flask, Google OAuth, YouTube API, pytest*

- Built a SaaS MVP that connects a creator’s YouTube account, tracks sponsored-post performance against their baseline, and generates shareable PDF reports for brands.

## EDUCATION

### Covenant University — B.Sc. Computer Science

2021 – 2025

Ota, Nigeria